


DeepSeek D=2 - allocated memory during one microbatch



- naive - peak 10.53 GiB (3 logit tensors [T, V] resident together)
- memory-efficient - peak 8.62 GiB (1 at a time)